

STAT 543/641 – WINTER 2019 – HOMEWORK #2

DUE MARCH ??, 2019

Let $\sigma(x) = (1 + e^{-x})^{-1}$ be the sigmoid function. The negative log-likelihood function associated to fitting a univariate logistic regression model to a dataset $(x_1, y_1), \dots, (x_n, y_n)$ is

$$\ell(a, b) = - \sum_{i=1}^n \left(y_i \log \sigma(a + bx_i) + (1 - y_i) \log (1 - \sigma(a + bx_i)) \right).$$

In this problem we will identify a condition under which $\ell(a, b)$ does not have a global minimum and a condition under which $\ell(a, b)$ has a unique local minimum.

(1) Let σ be the sigmoid function. Prove:

$$(*) \quad \sigma'(x) = \sigma(x)(1 - \sigma(x))$$

Conclude that $\sigma'(x) > 0$ for all x .

(2) Let x and y be constants and let

$$g(a, b) = y \log \sigma(a + bx) + (1 - y) \log (1 - \sigma(a + bx)).$$

Prove that

$$\nabla g(a, b) = (y - \sigma(a + bx)) \begin{bmatrix} 1 \\ x \end{bmatrix} \quad (\text{Hint: Use } (*).)$$

and that

$$\nabla^2 g(a, b) = -\sigma'(a + bx) \begin{bmatrix} 1 & x \\ x & x^2 \end{bmatrix}.$$

Show that $\nabla^2 g(a, b)$ is positive-semidefinite but not positive-definite, making $g(a, b)$ convex but not strictly convex.

(3) By (2) and the linearity of the Hessian operator ∇^2 ,

$$\nabla^2 \ell(a, b) = \sum_{i=1}^n Q_i(a, b), \quad \text{where} \quad Q_i(a, b) = \sigma'(a + bx_i) \begin{bmatrix} 1 & x_i \\ x_i & x_i^2 \end{bmatrix}$$

(4) Find a basis of the nullspace $N(Q_i(a, b))$ of $Q_i(a, b)$ whose elements do not depend on a and b .

(5) Suppose that there are indices i and j such that $x_i \neq x_j$. Prove that

$$\bigcap_{i=1}^n N(Q_i(a, b)) = \left\{ \begin{bmatrix} 0 \\ 0 \end{bmatrix} \right\}.$$

(Hint: Use (4).)

- (6) Suppose that there are indices i and j such that $x_i \neq x_j$. Show that $\nabla^2 \ell(a, b)$ is positive-definite and, hence, that $\ell(a, b)$ is strictly convex.
- (7) Conclude that if that there are indices i and j such that $x_i \neq x_j$, then maximum likelihood estimates for \hat{a} and \hat{b} are unique if they exist.

$$\nabla^2 \ell(a, b) = \sum_{i=1}^n Q_i(a, b), \quad \text{where} \quad Q_i(a, b) = \sigma'(a + bx_i) \begin{bmatrix} 1 & x_i \\ x_i & x_i^2 \end{bmatrix}$$

where

$$Q_i(a, b) = \sigma'(a + bx_i) \begin{bmatrix} 1 & x_i \\ x_i & x_i^2 \end{bmatrix}$$

$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

$$\begin{aligned} \frac{\partial g}{\partial a} &= y \frac{\sigma'(a + bx)}{\sigma(a + bx)} + (1 - y) \frac{(-\sigma'(a + bx))}{1 - \sigma(a + bx)} \\ &= y(1 - \sigma(a + bx)) - (1 - y)\sigma(a + bx) \\ &= y - y\sigma(a + bx) - \sigma(a + bx) + y\sigma(a + bx) \\ &= y - \sigma(a + bx) \end{aligned}$$

$$\begin{aligned} \frac{\partial g}{\partial b} &= y \frac{\sigma'(a + bx)x}{\sigma(a + bx)} + (1 - y) \frac{(-\sigma'(a + bx)x)}{1 - \sigma(a + bx)} \\ &= xy(1 - \sigma(a + bx)) - x(1 - y)\sigma(a + bx) \\ &= xy - xy\sigma(a + bx) - x\sigma(a + bx) + xy\sigma(a + bx) \\ &= x(y - \sigma(a + bx)) \end{aligned}$$

$$\begin{aligned} g''(b) &= -x^2 \sigma'(a + bx) \\ &= -x^2 \sigma(a + bx)(1 - \sigma(a + bx)) \end{aligned}$$

$$\begin{aligned} \ell(a, b) &= -\log L(a, b) = -\sum_{i=1}^n \left(y_i \log \sigma(a + bx_i) + (1 - y_i) \log (1 - \sigma(a + bx_i)) \right) \\ \ell''(a, b) &= \sum_{i=1}^n x_i^2 \sigma(a + bx_i)(1 - \sigma(a + bx_i)) \end{aligned}$$

For a vector $v \in \mathbb{R}^2$, write H_v for the open half plane

$$H_v = \{w \in \mathbb{R}^2 : v \cdot w > 0\}.$$

H_v is the connected component of $\mathbb{R}^2 - v^\perp$ containing v itself, v^\perp being the orthogonal complement of v :

$$v^\perp = \{w \in \mathbb{R}^2 : v \cdot w = 0\}$$

Write $C_{v,w}$ for the open cone spanned by vectors $v, w \in \mathbb{R}^2$, $v \neq w$:

$$C_{v,w} = \{av + bw : a, b > 0\}$$

- (1) Let $u, v, w \in \mathbb{R}^2$ be three distinct vectors. Prove that the following statements are equivalent:

- (a) $-u \in C_{v,w}$
- (b) $-u \in C_{v,w}$, $-w \in C_{u,v}$ and $-v \in C_{w,u}$
- (c) $H_u \cap H_v \cap H_w = \emptyset$
- (d) $H_u \cup H_v \cup H_w = \mathbb{R}^2$

If you're having trouble writing up a formal argument here, you can draw me a convincing, pretty diagram instead. (If you write up a proof, you can optionally include a picture!)

- (2) Let $x \in \mathbb{R}$. Show that the sets

$$\left\{ \begin{bmatrix} a \\ b \end{bmatrix} : \lim_{t \rightarrow \infty} \sigma(ta + tbx) = 0 \right\} \quad \text{and} \quad \left\{ \begin{bmatrix} a \\ b \end{bmatrix} : \lim_{t \rightarrow \infty} (1 - \sigma(ta + tbx)) = 0 \right\}$$

have the form $H_{\pm v}$, where $v = \begin{bmatrix} 1 \\ x \end{bmatrix}$.

- (3) For $K > 0$, identify the level curves

$$\left\{ \begin{bmatrix} a \\ b \end{bmatrix} : \sigma(a + bx) = K \right\} \quad \text{and} \quad \left\{ \begin{bmatrix} a \\ b \end{bmatrix} : (1 - \sigma(a + bx)) = K \right\}$$

as lines and the level sets

$$\left\{ \begin{bmatrix} a \\ b \end{bmatrix} : \sigma(a + bx) \leq K \right\} \quad \text{and} \quad \left\{ \begin{bmatrix} a \\ b \end{bmatrix} : (1 - \sigma(a + bx)) = K \right\}$$

as half-plane translates of the form $y + H_v$ and $z + H_w$.

- (4) Consider a three point data set $\{(x_1, y_1), (x_2, y_2), (x_3, y_3)\}$, where $x_1 < x_2 < x_3$ and let

$$L(a, b) = \prod_{i=1}^3 p(x_i, y_i | a, b),$$

where

$$p(x_i, y_i | a, b) = \sigma(a + bx_i)^{y_i} (1 - \sigma(a + bx_i)^{y_i})^{1-y_i}.$$

Show that

$$\lim_{a^2+b^2 \rightarrow \infty} L(a, b) = 0$$

if and only if $y_1 = y_3 = 1$ and $y_2 = 0$ or $y_1 = y_3 = 0$ and $y_2 = 0$. For $0 < K < 1$, consider the sets

$$S_{i,K} = \{(a, b) : p_i(x_i, y_i | a, b) > K\}, \quad i = 1, 2, 3.$$

Prove that $S_{1,K} \cap S_{2,K} \cap S_{3,K}$ is bounded if and only if either $y_1 = y_3 = 0$ and $y_2 = 1$ or $y_1 = y_3 = 1$ and $y_2 = 0$.

Let

$$v_i = \begin{bmatrix} 1 \\ x_i \end{bmatrix}, \quad i = 1, 2, 3.$$

Prove that

$$L(a, b)$$

(a) Suppose $y_1 = y_3 = 0$ and $y_2 = 1$.

Consider fitting a dataset $(x_1, y_1), \dots, (x_n, y_n)$ to the following simplified logistic regression model:

$$Y_i | X = x_i \sim p(x_i, y_i) := \sigma(bx_i)^{y_i} (1 - \sigma(bx_i))^{1-y_i}$$

Here, σ denotes the sigmoid function: $\sigma(t) = (1 + e^{-t})^{-1}$. Let $L(b)$ be the likelihood function associated this dataset/model.

Suppose that our dataset has the following property:

$$(*) \quad x_i < 0 \text{ if and only if } y_i = 0$$

Prove that

$$\lim_{b \rightarrow \infty} L(b) = 1.$$

What can we say about $\operatorname{argmax}_b L(b)$ in this case?

Solution:

Suppose that $x_i < 0$ and $y_i = 0$. Then $e^{-bx_i} \rightarrow \infty$ as $b \rightarrow \infty$. It follows that

$$\lim_{b \rightarrow \infty} p(x_i, y_i) = \lim_{b \rightarrow \infty} (1 - \sigma(bx_i)) = 1 - 0 = 1.$$

If $x_i > 0$ and $y_i = 1$, then $e^{-bx_i} \rightarrow 0$ as $b \rightarrow \infty$. It follows that

$$\lim_{b \rightarrow \infty} p(x_i, y_i) = \lim_{b \rightarrow \infty} \sigma(bx_i) = 1.$$

Therefore, by property (*),

$$\lim_{b \rightarrow \infty} p(x_i, y_i) = 1$$

for all $i = 1, \dots, n$. Consequently,

$$\lim_{b \rightarrow \infty} L(b) = \lim_{b \rightarrow \infty} \prod_{i=1}^n p(x_i, y_i) = 1.$$

Suppose, now, that there are indices i and j such that

$$x_i < 0 \text{ and } y_i = 1 \quad \text{and} \quad x_j > 0 \text{ and } y_j = 0.$$

In particular, property $(*)$ is not satisfied. Prove that

$$\lim_{b \rightarrow \pm\infty} L(b) = 0.$$

What can we say about $\operatorname{argmax}_b L(b)$ in this case?

Solution: As $y_i = 0$, $p(x_i, y_i|b) = \sigma(bx_i)$. As $x_i < 0$, $\sigma(bx_i) \rightarrow 0$ as $b \rightarrow \infty$. Thus

$$\lim_{b \rightarrow \infty} p(x_i, y_i|b) = \lim_{b \rightarrow \infty} \sigma(bx_i) = 0.$$

As $y_j = 1$, $p(x_j, y_j|b) = 1 - \sigma(bx_j)$. As $x_j > 0$, $1 - \sigma(bx_j) \rightarrow 0$ as $b \rightarrow \infty$. Thus

$$\lim_{b \rightarrow \infty} p(x_j, y_j|b) = \lim_{b \rightarrow \infty} \sigma(bx_j) = 0.$$

Let X_1, \dots, X_m and Y_1, \dots, Y_n be random samples from normally distributed populations with means μ_X and μ_Y and variances σ_X^2 and σ_Y^2 , respectively. Let S_X^2 and S_Y^2 be the standard unbiased estimators of σ_X^2 and σ_Y^2 , respectively.

(1) Suppose $\sigma_X^2 = \sigma_Y^2$ and write σ^2 for this common value.

$$S^2 := \frac{(m-1)S_X^2 + (n-1)S_Y^2}{m+n-2}$$

is an unbiased estimator of σ^2 . It's called the *pooled variance estimator*.

Solution: Since S_X^2 and S_Y^2 are unbiased estimators of σ_X and σ_Y , respectively,

$$\begin{aligned} \mathbb{E}[S^2] &= \frac{m-1}{m+n-2} \mathbb{E}[S_X^2] + \frac{n-1}{m+n-2} \mathbb{E}[S_Y^2] \\ &= \frac{m-1}{m+n-2} \sigma_X^2 + \frac{n-1}{m+n-2} \sigma_Y^2 \\ &= \frac{m-1}{m+n-2} \sigma^2 + \frac{n-1}{m+n-2} \sigma^2 && (\text{as } \sigma_X^2 = \sigma^2 = \sigma_Y^2) \\ &= \sigma^2 \end{aligned}$$

(2) Suppose, in addition to having common variance, that the X_i are independent of the Y_i . What is the distribution of

$$\frac{(m+n-2)S^2}{\sigma^2}?$$

What is the variance of S^2 ?

Solution: Since S_X^2 and S_Y^2 are independent,

$$\frac{(m-1)S_X^2}{\sigma^2} \sim \chi_{m-1}^2 \quad \text{and} \quad \frac{(n-1)S_Y^2}{\sigma^2} \sim \chi_{n-1}^2,$$

and it follows from general properties of χ^2 -distributions that

$$\frac{(m+n-2)S^2}{\sigma^2} = \frac{1}{\sigma^2} ((m-1)S_X^2 + (n-1)S_Y^2) \sim \chi_{m+n-2}^2.$$

Since χ_{m+n-2}^2 -distributed random variable has variance $2(m+n-2)$, it follows that

$$\text{Var } S^2 = \frac{2\sigma^4}{m+n-2}.$$

(3*) Generalize these results from the case of $K = 2$ populations to that of an arbitrary K . Compare with equation (4.15) in [2].

(4*) Can you prove analogous results with covariance matrices in place of scalar variances?

Since S_X^2 and S_Y^2 are independent,

$$\frac{(m-1)S_X^2}{\sigma^2} \sim \chi_{m-1}^2 \quad \text{and} \quad \frac{(n-1)S_Y^2}{\sigma^2} \sim \chi_{n-1}^2,$$

and it follows from general properties of χ^2 -distributions that

$$\frac{(m+n-2)S^2}{\sigma^2} = \frac{1}{\sigma^2} ((m-1)S_X^2 + (n-1)S_Y^2) \sim \chi_{m+n-2}^2.$$

Since χ_{m+n-2}^2 -distributed random variable has variance $2(m+n-2)$, it follows that

$$\text{Var } S^2 = \frac{2\sigma^4}{m+n-2}.$$

[1, Exercise 12.16] This exercise examines an extreme case in which the likelihood equations for logistic regression have no solution.

Consider the following 20-point data set:

$$(0, 1), (0, 1), (0, 1), (0, 1), (0, 1), (0, 1), (0, 1), (0, 1), (0, 1), (0, 1) \\ (1, 0), (1, 0), (1, 0), (1, 0), (1, 0), (1, 1), (1, 1), (1, 1), (1, 1), (1, 1)$$

- (1) Observe that, empirically, $\text{Prob}(Y = 1|X = 0) = 1$ and $\text{Prob}(Y = 1|X = 1) = 0.5$. Let $\sigma(t) = (1 + e^{-t})^{-1}$ be the sigmoid function. Are there a and b such that $\sigma(a + b \cdot 0) = 1$ or $\sigma(a + b \cdot 1) = 0.5$?

Solution:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Since the exponential function takes values in $(0, \infty)$,

$$0 < \sigma(x) < 1,$$

for every $x \in \mathbb{R}$. In particular, there are no numbers a and b such that $\sigma(a + b \cdot 0) = 1$. Clearly, $\sigma(0) = \sigma(0 + 0 \cdot 1) = 0.5$.

- (2) Let $\mathcal{L}(a, b)$ be the likelihood function associated to fitting a logistic regression model to this data set. Show that

$$L := \lim_{b \rightarrow \infty} \mathcal{L}(-b, b) = \sup_{(a, b) \in \mathbb{R}^2} \mathcal{L}(a, b) < \infty$$

and that $\mathcal{L}(a, b) \neq L$ for any $(a, b) \in \mathbb{R}^2$. What are

$$\lim_{b \rightarrow \infty} \sigma(-b + b \cdot 0) \quad \text{and} \quad \lim_{b \rightarrow \infty} \sigma(-b + b \cdot 1)?$$

Let (\mathbf{X}, Y) be jointly distributed, where \mathbf{X} is a p -dimensional random vector and Y takes values in $\{1, \dots, K\}$. Suppose that, for each k , $\mathbf{X}|Y = k$ has Gaussian distribution with mean $\boldsymbol{\mu}_k$ and variance Σ , with the latter independent of k .

Consider a data set $(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$, where $\mathbf{x}^{(i)} \in \mathbb{R}^{1 \times p}$ and $y^{(i)} \in \{1, \dots, K\}$. For $1 \leq k \leq K$, let

$$I_k = \{i : y^{(i)} = k\}, \quad n_k = |I_k|, \quad \hat{\pi}_k = \frac{n_k}{n}.$$

Define sample means $\boldsymbol{\mu}_k$ and a pooled sample covariance Σ by

$$\begin{aligned} \hat{\boldsymbol{\mu}}_k &= \hat{\boldsymbol{\mu}}_{k, \mathbf{x}} = \frac{1}{n_k} \sum_{i \in I_k} \mathbf{x}^{(i)} \in \mathbb{R}^{p \times 1}, \\ \hat{\Sigma} &= \hat{\Sigma}_{\mathbf{x}} = \frac{1}{n - K} \sum_{k=1}^K \sum_{i \in I_k} (\mathbf{x}^{(i)} - \hat{\boldsymbol{\mu}}_k)^T (\mathbf{x}^{(i)} - \hat{\boldsymbol{\mu}}_k) \in \mathbb{R}^{p \times p}. \end{aligned}$$

Define linear discriminant functions, $\delta_k = \delta_{k, \mathbf{x}}$, by

$$\delta_k(\mathbf{v}) = \delta_{k, \mathbf{x}}(\mathbf{v}) = \mathbf{v} \hat{\Sigma} \hat{\boldsymbol{\mu}}_k^T - \frac{1}{2} \hat{\boldsymbol{\mu}}_k \hat{\Sigma} \hat{\boldsymbol{\mu}}_k^T + \log \hat{\pi}_k, \quad \mathbf{v} \in \mathbb{R}^{p \times 1}.$$

Let $\mathbf{a} \in \mathbb{R}^{p \times 1}$ and let

$$\begin{aligned} \mathbf{w}^{(i)} &= \mathbf{x}^{(i)} - \mathbf{a}. \\ \hat{\boldsymbol{\mu}}_{k, \mathbf{w}} &= \hat{\boldsymbol{\mu}}_{k, \mathbf{x}} - \mathbf{a}, \quad \Sigma_{\mathbf{w}} = \Sigma_{\mathbf{x}} \\ \delta_{k_1, \mathbf{w}}(v - a) - \delta_{k_2, \mathbf{w}}(v - a) &= \delta_{k_1, \mathbf{x}}(v) - \delta_{k_2, \mathbf{x}}(v) \end{aligned}$$

Let $U \in \mathbb{R}^{p \times p}$ be an orthogonal matrix and let $w^{(i)} = Ux^{(i)}$. Then

$$\delta_{k, Ux}(Uv) = \delta_x(v).$$

$$\sum_{k=1}^K \pi_k \mu_k = \mu$$

Suppose

$$\begin{aligned} Y &\sim \text{Categorical}\left(\frac{1}{K}, \dots, \frac{1}{K}\right), \\ \mathbf{X}|Y = k &\sim N(\boldsymbol{\mu}_k, \sigma I). \end{aligned}$$

Show that

$$\operatorname{argmin}_k p(Y = k | \mathbf{X} = \mathbf{x}) = \operatorname{argmin}_k \|\mathbf{x} - \boldsymbol{\mu}_k\|.$$

Logistic regression (with and without ridge regularization, with and without PCA), LDA, Gaussian naïve Bayes, for breast cancer data set. Plot in 2d with decision boundary. Optional: Lasso

Document classification with multinomial naïve Bayes

Ridge regression via constrained optimization.

1. TOTAL VARIATION = VARIANCE WITHIN + VARIANCE BETWEEN

Let X and Y be jointly distributed random variables, where

$$\begin{aligned} Y &\sim \text{Categorical}(\pi_1, \dots, \pi_K), & \sum \pi_k &= 1, \\ X | Y = k &\sim N(\mu_k, \sigma_k^2), & k &= 1, \dots, K. \end{aligned}$$

Assume that the μ_k are pairwise distinct. Prove that

$$\mathbb{E}[X] = \sum_{k=1}^K \pi_k \mu_k$$

and that

$$\text{Var } X = \sum_{k=1}^K \pi_k \sigma_k^2 + \sum_{k=1}^K \pi_k (\mu_k - \mu)^2.$$

To establish the decomposition of the variance, you might want to use the *law of total variation*:

$$\text{Var } X = \mathbb{E}(\text{Var}(X|Y)) + \text{Var } \mathbb{E}(X|Y).$$

The random variable X has marginal density $p(x)$, where

$$p(x) = \sum_{k=1}^K p(x, k) = \sum_{k=1}^K p(k) p(x|k) = \sum_{k=1}^K \pi_k p(x|k).$$

Therefore,

$$\mathbb{E}[X] = \sum_{k=1}^K \pi_k \mathbb{E}[X_k]$$

where X_k is a random variable with density $p(x|k)$. By hypothesis, X_k has expected value μ_k . Therefore,

$$\mathbb{E}[X] = \sum_{k=1}^K \pi_k \mu_k.$$

Since $\text{Var}(X|Y = k) = \sigma_k^2$,

$$\mathbb{E}[\text{Var}(X|Y)] = \sum_{k=1}^K p(Y = k) \text{Var}(X|Y = k) = \sum_{k=1}^K \pi_k \sigma_k^2.$$

Let $\mu = \mathbb{E}[X]$. By the *law of total expectation*,

$$\mathbb{E}[\mathbb{E}[X|Y]] = \mathbb{E}[X] = \mu.$$

Let $Z = \mathbb{E}[X|Y]$. Then set of values of Z is $\{\mu_1, \dots, \mu_K\}$. Since the μ_k are pairwise distinct,

$$\text{Prob}[Z = \mu_k] = p(Y = k) = \pi_k.$$

Therefore,

$$\text{Var } Z = \sum_{k=1}^K \text{Prob}[Z = \mu_k] (\mu_k - \mu)^2 = \sum_{k=1}^K \pi_k (\mu_k - \mu)^2.$$

Let x_1, \dots, x_n be numbers and set

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i, \quad \sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2.$$

Suppose that

$$\{1, \dots, n\} = \bigsqcup_{k=1}^K I_k \quad (\text{disjoint union})$$

with

$$n_k := |I_k| > 0.$$

Set

$$\mu_k = \frac{1}{n_k} \sum_{i \in I_k} x_i, \quad \sigma_k^2 = \frac{1}{n_k} \sum_{i \in I_k} (x_i - \mu_k)^2$$

Prove that

$$\sigma^2 = \sum_{k=1}^K \pi_k \sigma_k^2 + \sum_{k=1}^K \pi_k (\mu_k - \mu)^2, \quad \text{where } \pi_k = \frac{n_k}{n}.$$

(This is an “algebraic version” of the law of total variance)

$$\begin{aligned}
\frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2 &= \frac{1}{n} \sum_{k=1}^K \sum_{i \in I_k} (x_i - \mu)^2 \\
&= \frac{1}{n} \sum_{k=1}^K \sum_{i \in I_k} (x_i - \mu_k + \mu_k - \mu)^2 \\
&= \frac{1}{n} \sum_{k=1}^K \sum_{i \in I_k} \{ (x_i - \mu_k)^2 + (\mu_k - \mu)^2 + 2(x_i - \mu_k)(\mu_k - \mu) \} \\
&= \frac{1}{n} \sum_{k=1}^K n_k \frac{1}{n_k} \sum_{i \in I_k} (x_i - \mu_k)^2 + \frac{1}{n} \sum_{k=1}^K (\mu_k - \mu)^2 \sum_{i \in I_k} 1 \\
&\quad + \frac{2}{n} \sum_{k=1}^K (\mu_k - \mu) \sum_{i \in I_k} (x_i - \mu_k) \\
&= \sum_{k=1}^K \frac{n_k}{n} \sigma_k^2 + \sum_{k=1}^K \frac{n_k}{n} (\mu_k - \mu)^2 + 0 \\
&= \sum_{k=1}^K \pi_k \sigma_k^2 + \sum_{k=1}^K \pi_k (\mu_k - \mu)^2
\end{aligned}$$

Define a probability space (Ω, μ) by

$$\Omega = \{x_1, \dots, x_n\} \times \{1, \dots, K\}, \quad \mu(\{(x_i, k)\}) = p_{ik} := \begin{cases} \frac{1}{n} & \text{if } i \in I_k, \\ 0 & \text{otherwise.} \end{cases}$$

Note that μ is, indeed, a probability measure on Ω :

$$\sum_{(i,k) \in \Omega} p_{ik} = \sum_{k=1}^K \sum_{i=1}^n p_{ik} = \sum_{k=1}^K \sum_{i \in I_k} \frac{1}{n} = \sum_{k=1}^K \frac{n_k}{n} = 1$$

Define random variables X and Y on Ω by

$$X(x_i, k) = x_i \quad \text{and} \quad Y(x_i, k) = k.$$

Then

$$\mu = \mathbb{E}[X] = \sum_{k=1}^K \sum_{i=1}^n p_{ik} x_i = \sum_{k=1}^K \sum_{i \in I_k} \frac{1}{n} x_i = \frac{1}{n} \sum_{k=1}^K \sum_{i \in I_k} x_i = \frac{1}{n} \sum_{i=1}^n x_i.$$

Let $Z = \mathbb{E}[X|Y]$, so that

$$Z(k) = \mathbb{E}[X|Y = k] = \frac{\sum_{i=1}^n p_{ik} x_i}{\sum_{i=1}^n p_{ik}} = \frac{\sum_{i \in I_k} x_i}{\sum_{i \in I_k} 1} = \frac{1}{n_k} \sum_{i \in I_k} x_i = \mu_k.$$

$$\begin{aligned}
\text{Var}[X|Y = k] &= \mathbb{E}[(X - Z)^2|Y = k] \\
&= \frac{\sum_{i=1}^n p_{ik}(X(i) - Z(k))^2}{\sum_{i=1}^n p_{ik}} \\
&= \frac{\frac{1}{n} \sum_{i \in I_k} (x_i - \mu_k)^2}{\sum_{i \in I_k} \frac{1}{n}} \\
&= \frac{1}{n_k} \sum_{i \in I_k} (x_i - \mu_k)^2 \\
&= \sigma_k^2
\end{aligned}$$

$$\begin{aligned}
\mathbb{E}[\text{Var}[X|Y]] &= \sum_{k=1}^K \text{Prob}(Y = k) \text{Var}[X|Y = k] \\
&= \sum_{k=1}^K \left(\sum_{(i,k)} p_{ik} \right) \sigma_k^2 \\
&= \sum_{k=1}^K \frac{n_k}{n} \sigma_k^2 \\
&= \sum_{k=1}^K \pi_k \sigma_k^2
\end{aligned}$$

Finally,

$$\begin{aligned}
\text{Var } Z &= \mathbb{E}[(Z - \mathbb{E}[Z])^2] \\
&= \mathbb{E}[(Z - \mu)^2] && \text{(by the law of total expectation)} \\
&= \sum_{k=1}^K \text{Prob}(Z = \mu_k) (\mu_k - \mu)^2 \\
&= \sum_{k=1}^K \text{Prob}(Y = k) (\mu_k - \mu)^2 && \text{(as } k = \ell \iff \mu_k = \mu_\ell) \\
&= \sum_{k=1}^K \left(\sum_{(i,k)} p_{ik} \right) (\mu_k - \mu)^2 \\
&= \sum_{k=1}^K \pi_k (\mu_k - \mu)^2.
\end{aligned}$$

1.1. Matrix version.

$$\begin{aligned}
\frac{1}{n} \sum_{i=1}^n (x_i - \mu)(x_i - \mu)^T &= \frac{1}{n} \sum_{k=1}^K \sum_{i \in I_k} (x_i - \mu)(x_i - \mu)^T \\
&= \frac{1}{n} \sum_{k=1}^K \sum_{i \in I_k} ((x_i - \mu_k) + (\mu_k - \mu))((x_i - \mu_k) + (\mu_k - \mu))^T \\
&= \frac{1}{n} \sum_{k=1}^K \sum_{i \in I_k} \{ (x_i - \mu_k)(x_i - \mu_k)^T + (\mu_k - \mu)(\mu_k - \mu)^T \\
&\quad + (x_i - \mu_k)(\mu_k - \mu)^T + (\mu_k - \mu)(x_i - \mu_k)^T \} \\
&= \frac{1}{n} \sum_{k=1}^K n_k \frac{1}{n_k} \sum_{i \in I_k} (x_i - \mu_k)(x_i - \mu_k)^T + \frac{1}{n} \sum_{k=1}^K (\mu_k - \mu)(\mu_k - \mu)^T \sum_{i \in I_k} 1 \\
&\quad + \frac{1}{n} \sum_{k=1}^K \left\{ \sum_{i \in I_k} (x_i - \mu_k) \right\} (\mu_k - \mu)^T \\
&\quad + \frac{1}{n} \sum_{k=1}^K (\mu_k - \mu) \sum_{i \in I_k} (x_i - \mu_k)^T \\
&= \sum_{k=1}^K \frac{n_k}{n} \Sigma_k^2 + \sum_{k=1}^K \frac{n_k}{n} (\mu_k - \mu)(\mu_k - \mu)^T + 0 + 0 \\
&= \sum_{k=1}^K \pi_k \Sigma_k^2 + \sum_{k=1}^K \pi_k (\mu_k - \mu)(\mu_k - \mu)^T
\end{aligned}$$

REFERENCES

- [1] Casella, Berger, *Statistical Inference (2nd ed.)*, Duxbury, 2002.
- [2]